

AI Client Rollout SOP

Purpose

This SOP defines:

- the recommended initial AI settings baseline
- the runtime-control baseline
- the rollout sequence to use for each client
- the launch gate and weekly review checklist

Use this SOP for every new client install unless there is a documented exception.

Default Baseline Settings

AI governance

- AI context strictness: strict
- Lock AI mode: auto
- Minimum confidence for AI advice: 0.90
- Minimum confidence for AI actions: 0.94
- Minimum goal relevance for proactive guidance: 0.72
- Allow AI auto-complete tasks: enabled
- Auto-complete confidence threshold: 0.97
- Enable autonomous AI threshold tuning: enabled

Incident alerting

- Enable AI incident alerts: enabled
- Enable AI incident checks: enabled
- Medium cooldown: 360 minutes
- High cooldown: 240 minutes
- Critical cooldown: 120 minutes

Runtime Control Baseline

Set these in the AI Control Center at first deployment:

- Global: normal
- Coach: normal
- Clarity Chat: normal
- Assistant: suggest_only
- Customer Thread: suggest_only
- Commercial Assistant: suggest_only
- Task Automation: normal
- Autonomous Tuning: diagnostics_only

Client Rollout Tiers

Tier A: Early-stage / low-process client

Use when:

- low data quality
- few users
- no dedicated operator
- incomplete products, pipeline, or invoicing setup

Starting posture:

- Assistant: suggest_only
- Customer Thread: suggest_only
- Commercial Assistant: suggest_only
- Autonomous Tuning: diagnostics_only

Tier B: Operationally mature client

Use when:

- good deal/contact hygiene
- invoicing configured
- operator available
- scheduler and alerts monitored

Starting posture:

- Assistant: suggest_only
- Customer Thread: suggest_only
- Commercial Assistant: suggest_only
- Autonomous Tuning: diagnostics_only

Promotion to normal can happen earlier than Tier A after review.

Tier C: High-volume / high-risk client

Use when:

- larger team
- higher message volume
- revenue-impacting automation is active
- rapid issue propagation is possible

Starting posture:

- Assistant: suggest_only
- Customer Thread: suggest_only
- Commercial Assistant: suggest_only
- Autonomous Tuning: diagnostics_only
- Keep stricter review cadence and do not accelerate rollout without explicit signoff.

Standard Rollout Sequence

Stage 0: Internal validation

Required before client-visible rollout:

- migrations complete
- all scheduler jobs running

- incident alerts tested
- prompt rollback tested
- AI Control Center accessible
- diagnostics workbench accessible

Keep:

- Assistant: suggest_only
- Customer Thread: suggest_only
- Commercial Assistant: suggest_only
- Autonomous Tuning: diagnostics_only

Stage 1: Guided usage

Enable:

- Coach
- Clarity Chat
- Task Automation

Keep:

- Assistant: suggest_only
- Customer Thread: suggest_only
- Commercial Assistant: suggest_only

Daily review:

- AI incidents
- blocked spikes
- prompt warnings
- stale jobs
- approval spikes

Minimum duration:

- 5 business days

Stage 2: Limited execution

Move to:

- Assistant: normal

Keep:

- Customer Thread: suggest_only
- Commercial Assistant: suggest_only
- Autonomous Tuning: diagnostics_only unless review is clean

Promotion requirements:

- no unresolved high AI incidents
- no prompt regression high risk
- no stale calibration or reconciliation jobs

- low blocked spike rate

Minimum duration:

- 5 business days

Stage 3: Controlled client-facing execution

Move to:

- Customer Thread: normal

Keep Commercial Assistant on suggest_only until:

- invoicing is configured
- contact/deal linkage is reliable
- approval workflow is active
- rejection and reversal rates are acceptable

Minimum duration:

- 5 business days

Stage 4: Commercial execution

Move to:

- Commercial Assistant: normal
- Autonomous Tuning: normal

Only if:

- no unresolved high incidents
- approval rejection rate is acceptable
- prompt regression is not active
- operator is actively monitoring incidents and diagnostics

Hard Promotion Rules

Do not broaden a client rollout unless all of the following are true:

- incident alerts are configured and tested
- ai_incident_check.php is running every 15 minutes
- ai_outcome_reconciliation.php is healthy
- ai_confidence_calibration.php is healthy
- workflow_scheduler.php is healthy
- ai_runtime_control_cleanup.php is healthy
- no unresolved high or critical AI incidents
- prompt rollback has been tested
- diagnostics and control center are accessible to the operator

Launch Gate Checklist

Before launch for any client, confirm:

- company profile completed
- products and pricing configured
- invoicing configured if commercial assistant will be used
- email/inbox working
- permissions and admin access verified
- prompt control page accessible
- diagnostics page accessible
- incident workbench accessible
- alert channels validated
- runtime controls tested

Weekly Review Checklist

Review once per week per active client:

- active incidents count
- blocked decision spikes
- approval-required spikes
- prompt regression warnings
- stale-context and overload rates
- scheduler health
- autonomous tuning changes in last 7 days
- fallback usage trends
- commercial approval rejection rate

Safe Mode Policy

If a client experiences repeated degradation:

Set immediately:

- Assistant: suggest_only
- Customer Thread: suggest_only
- Commercial Assistant: suggest_only
- Autonomous Tuning: diagnostics_only

If issues persist:

- pause affected surface
- keep diagnostics active
- investigate incidents, prompt regression, job health, and prompt rollback options

Recommended Default Profile

Use this as the default profile for most clients:

- AI context strictness: strict
- AI mode lock: auto
- Advice confidence: 0.90
- Action confidence: 0.94
- Goal relevance: 0.72
- Auto task completion confidence: 0.97
- Assistant: suggest_only

- Customer Thread: suggest_only
- Commercial Assistant: suggest_only
- Autonomous Tuning: diagnostics_only

This is the default launch-safe baseline.